

Frequent Item Sets

Chau Tran & Chun-Che Wang

Outline

1. Definitions
 - Frequent Itemsets
 - Association rules
2. Apriori Algorithm

Frequent Itemsets

What? Why? How?

Motivation 1: Amazon suggestions

Frequently Bought Together



+



+



Price for all three: **\$85.90**



Add all three to Cart

Add all three to Wish List

[Show availability and shipping details](#)

- ✓ **This item:** Kit Kat Candy Bar, Crisp Wafers in Milk Chocolate, 1.5-Ounce Bars (Pack of 36) **\$28.63** (\$0.53 / oz)
- ✓ Reese's Peanut Butter Cups, 1.5-Ounce Packages (Pack of 36) **\$24.30** (\$0.45 / oz)
- ✓ Twix-chocolate Caramel Cookie Bars, 36ct **\$32.97** (\$9.16 / 10 Items)

Amazon suggestions (German version)

ALUMINIUM Baseballschläger 30" American Baseball
von [Outdoor 4 You - Shop](#)
★★★★☆ (4 Kundenrezensionen) [Mehr zu diesem Artikel](#)

Preis: **EUR 17,58**

Auf Lager.
Verkauf und Versand durch [NORMANI](#).
Noch 5 Stück auf Lager.
[4 neu](#) ab EUR 17,58

Marken-Uhren mit Tiefpreis-Garantie finden Sie im [Uhren-Shop](#) bei Amazon.de/Uhren.

Kunden, die diesen Artikel gekauft haben, kauften auch Seite 1 von 23

			
Leder Quarzaarbeitshandschuhe schwarz S-XXL	Balaclava 3-Loch ★★★★☆ (4) EUR 3,50	Pfefferspray KO-FOG 40ML ★★★★★ (9) EUR 5,05	Baseballschläger Holz 32" American Baseball natur

Motivation 2: Plagiarism detector

- Given a set of documents (eg. homework handin)
 - Find the documents that are similar

Motivation 3: Biomarker

- Given the set of medical data
 - For each patient, we have his/her genes, blood proteins, diseases
 - Find patterns
 - which genes/proteins cause which diseases

What do they have in common?

- A large set of **items**
 - things sold on Amazon
 - set of documents
 - genes or blood proteins or diseases
- A large set of **baskets**
 - shopping carts/orders on Amazon
 - set of sentences
 - medical data for multiple of patients

Goal

- Find a general many-many mapping between two set of items
 - $\{\text{Kitkat}\} \Rightarrow \{\text{Reese, Twix}\}$
 - $\{\text{Document 1}\} \Rightarrow \{\text{Document 2, Document 3}\}$
 - $\{\text{Gene A, Protein B}\} \Rightarrow \{\text{Disease C}\}$

Approach

- $A = \{A_1, A_2, \dots, A_m\}$
- $B = \{B_1, B_2, \dots, B_n\}$

A, B are subset of
 I = set of items

$$\begin{aligned} P(B|A) &= \frac{P(A,B)}{P(A)} \\ &= \frac{\text{Count}(A,B)}{\text{Count}(A)} \end{aligned}$$

Definitions

- **Support** for itemset A: Number of baskets containing all items in A
 - Same as $\text{Count}(A)$
- Given a **support threshold** s , the set of items that appear in at least s baskets are called **frequent itemsets**

Example: Frequent Itemsets

- Items = {milk, coke, pepsi, beer, juice}

B1 = {m,c,b}	B2 = {m,p,j}
B3 = {m,b}	B4 = {c,j}
B5 = {m, p, b}	B6 = {m,c,b,j}
B7 = {c,b,j}	B8 = {b,c}

- Frequent itemsets for support threshold = 3:
 - {m}, {c}, {b}, {j}, {m,b}, {b,c}, {c,j}

Association Rules

- $A \Rightarrow B$ means: “if a basket contains items in A, it is likely to contain items in B”
- There are exponentially many rules, we want to find **significant/interesting** ones
- Confidence of an association rule:
 - $\text{Conf}(A \Rightarrow B) = P(B \mid A)$

Interesting association rules

- Not all high-confidence rules are interesting
 - The rule $X \Rightarrow \text{milk}$ may have high confidence for many itemsets X , because milk is just purchased very often (independent of X), and the confidence will be high
- Interest of an association rule:
 - $\text{Interest}(A \Rightarrow B) = \text{Conf}(A \Rightarrow B) - P(B)$
 $= P(B \mid A) - P(B)$

- $\text{Interest}(A \Rightarrow B) = P(B \mid A) - P(B)$
 - > 0 if $P(B \mid A) > P(B)$
 - $= 0$ if $P(B \mid A) = P(B)$
 - < 0 if $P(B \mid A) < P(B)$

Example: Confidence and Interest

B1 = {m,c,b}	B2 = {m,p,j}
B3 = {m,b}	B4 = {c,j}
B5 = {m, p, b}	B6 = {m,c,b,j}
B7 = {c,b,j}	B8 = {b,c}

- Association rule: $\{m,b\} \Rightarrow c$
 - Confidence = $2/4 = 0.5$
 - Interest = $0.5 - \frac{5}{8} = -\frac{1}{8}$
 - High confidence but not very interesting

Overview of Algorithm

- Step 1: Find all frequent itemsets I
- Step 2: Rule generation
 - For every subset A of I , generate a rule $A \Rightarrow I \setminus A$
 - Since I is frequent, A is also frequent
 - Output the rules above the confidence threshold

Example: Finding association rules

B1 = {m,c,b}	B2 = {m,p,j}
B3 = {m,b}	B4 = {c,j}
B5 = {m, p, b}	B6 = {m,c,b,j}
B7 = {c,b,j}	B8 = {b,c}

- Min support $s=3$, confidence $c=0.75$

- 1) Frequent itemsets:

- {b,m} {b,c} {c,n} {c,j} {m,c,b}

- 2) Generate rules:

- ~~$b \Rightarrow m = 4/6$~~ $b \Rightarrow c = 5/6$

- $m \Rightarrow b = 4/5$...

- $b,m \Rightarrow c = 3/4$

- ~~$b,c \Rightarrow m = 3/5$~~

...

How to find frequent itemsets?

- Have to find subsets A such that $\text{Support}(A) > s$
 - There are 2^n subsets
 - Can't be stored in memory

How to find frequent itemsets?

- Solution: only find subsets of size 2



Really?


- Frequent pairs are common, frequent triples are rare, don't even talk about $n=4$
- Let's first concentrate on pairs, then extend to larger sets (wink at Chun)
- The approach
 - Find $\text{Support}(A)$ for all A such that $|A| = 2$

Naive Algorithm


- For each basket b :
 - for each pair (i_1, i_2) in b :
 - increment count of (b_1, b_2)
- Still fail if $(\#items)^2$ exceeds main memory
 - Walmart has 10^5 items
 - Counts are 4-byte integers
 - Number of pairs = $10^5 * (10^5 - 1) / 2 = 5 * 10^9$
 - $2 * 10^{10}$ bytes (20 GB) of memory needed

Not all pairs are equal

- Store a hash table
 - $(i1, i2) \Rightarrow \text{index}$
- Store triples $[i1, i2, c(i1, i2)]$
 - uses 12 bytes per pair
 - but only for pairs with count > 0
- Better if less than $\frac{1}{3}$ of possible pairs actually occur



Triangular Matrix



Triples

Summary

- What?
 - Given a large set of **baskets** of **items**, find items that are correlated
- Why?
- How?
 - Find frequent itemsets
 - subsets that occur more than s times
 - Find association rules
 - $\text{Conf}(A \Rightarrow B) = \text{Support}(A, B) / \text{Support}(A)$

A-Priori Algorithm

Naive Algorithm Revisited



- Pros:
 - Read the entire file (transaction DB) once
- Cons
 - Fail if $(\text{\#items})^2$ exceeds main memory

A-Priori Algorithm

- Designed to reduce the number of pairs that need to be counted
- How?
hint: There is no such thing as a free lunch
- Perform 2 passes over data

A-Priori Algorithm

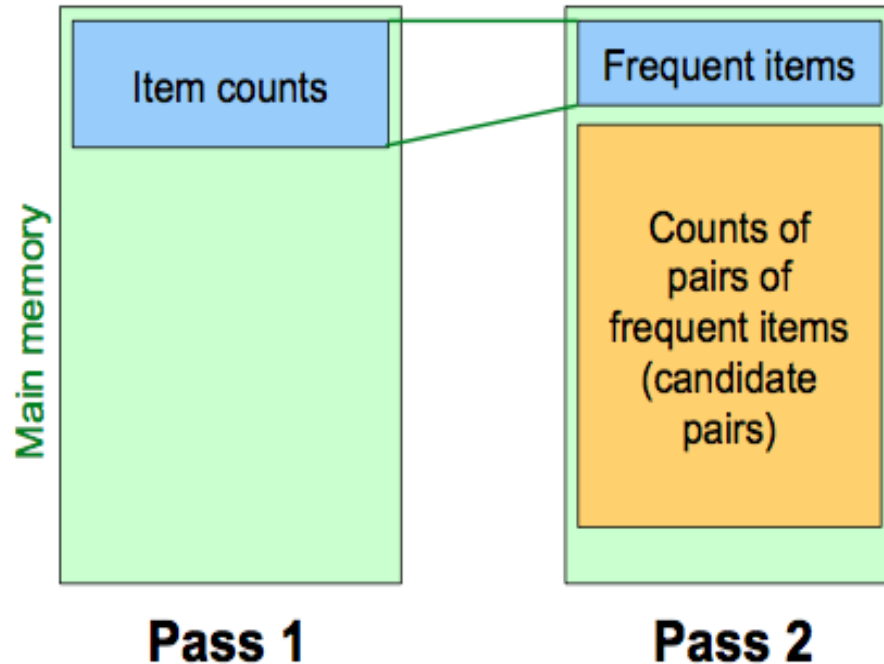
- **Key idea : monotonicity**
 - If a set of items appears at least s times, so does every subset
- **Contrapositive for pairs**
 - If item i does not appear in s baskets, then no pair including i can appear in s baskets



A-Priori Algorithm

- Pass 1:
 - Count the occurrences of each **individual item**
 - items that appear at least s time are the frequent items
- Pass 2:
 - Read baskets again and count in only those pairs where both elements are frequent (from pass 1)

A-Priori Algorithm

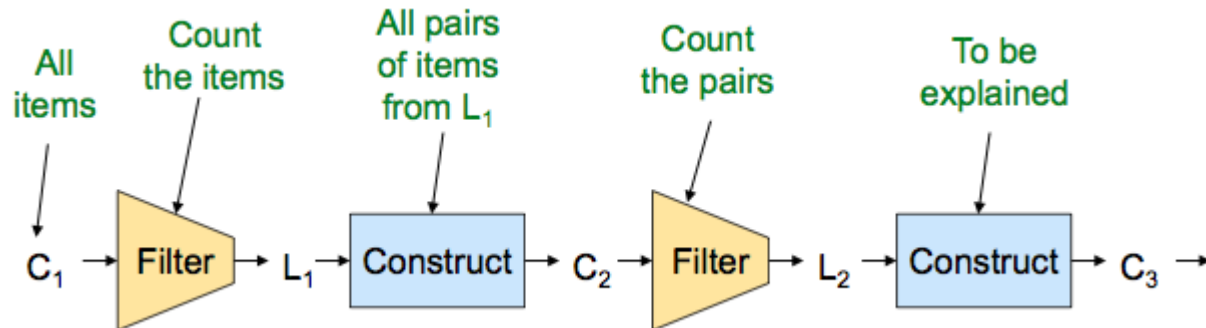


Frequent Triples, Etc.

For each k , we construct two sets of k -tuples

C_k Candidate k -tuples = those might be frequent sets (support $> s$)

L_k The set of truly frequent k -tuples



Example

■ Hypothetical steps of the A-Priori algorithm

- $C_1 = \{ \{b\} \{c\} \{j\} \{m\} \{n\} \{p\} \}$
- Count the support of itemsets in C_1
- Prune non-frequent: $L_1 = \{ b, c, j, m \}$
- Generate $C_2 = \{ \{b,c\} \{b,j\} \{b,m\} \{c,j\} \{c,m\} \{j,m\} \}$
- Count the support of itemsets in C_2
- Prune non-frequent: $L_2 = \{ \{b,m\} \{b,c\} \{c,m\} \{c,j\} \}$
- Generate $C_3 = \{ \{b,c,m\} \{b,c,j\} \{b,m,j\} \{c,m,j\} \}$
- Count the support of itemsets in C_3
- Prune non-frequent: $L_3 = \{ \{b,c,m\} \}$

A-priori for All Frequent Itemsets

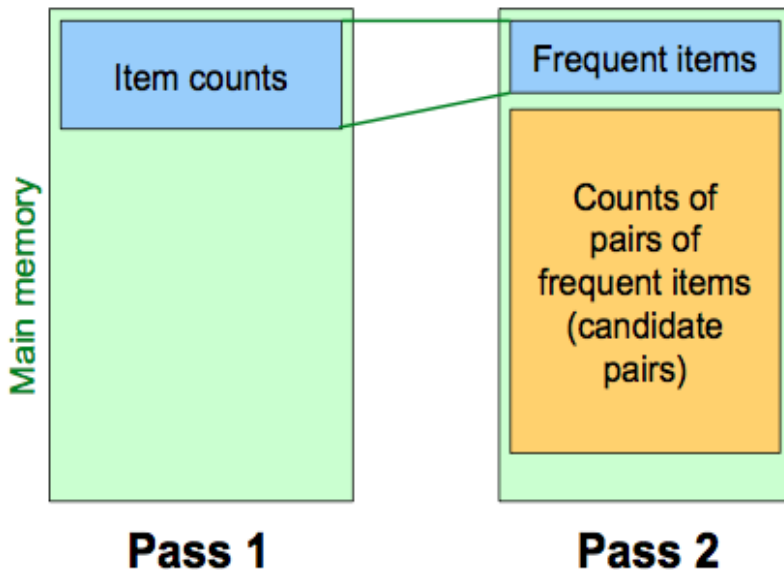
- For finding frequent k-tuple: Scan entire data k times
- Needs room in main memory to count each candidate k-tuple
- Typical, $k = 2$ requires the most memory

What else can we improve?

- Observation

In pass 1 of a-priori, most memory is idle !

Can we use the idle memory to reduce memory required in pass 2?



PCY Algorithm

- PCY (Park-Chen-Yu) Algorithm
- Take advantage of the idle memory in pass1
 - During pass 1, maintain a hash table
 - Keep a count for each bucket into which pairs of items are hashed

```
FOR (each basket) :  
    FOR (each item in the basket) :  
        add 1 to item's count;  
    New  
in  
PCY { FOR (each pair of items) :  
        hash the pair to a bucket;  
        add 1 to the count for that bucket;
```

PCY Algorithm - Pass 1

New in PCY

```
FOR (each basket) :  
  FOR (each item in the basket) :  
    add 1 to item's count;  
  FOR (each pair of items) :  
    hash the pair to a bucket;  
    add 1 to the count for that bucket;
```

Define the hash function: $h(i, j) = (i + j) \% 5 = K$
(Hashing pair (i, j) to bucket K)

Pass 1

Items A = {milk, coke, cookies, pepsi, juice},. Baskets are

$B_1 = \{m, c, b\}$ $B_2 = \{m, p, j\}$
 $B_3 = \{m, b\}$ $B_4 = \{c, j\}$
 $B_5 = \{m, p, b\}$ $B_6 = \{m, c, b, j\}$
 $B_7 = \{c, b, j\}$ $B_8 = \{b, c\}$

Support threshold $s=3$,

If we assign milk=1, coke=2, cookies= 3, pepsi=4, and juice=5
Then,

$B_1 = \{1, 2, 3\}$ $B_2 = \{1, 4, 5\}$
 $B_3 = \{1, 3\}$ $B_4 = \{2, 5\}$
 $B_5 = \{1, 3, 4\}$ $B_6 = \{1, 2, 3, 5\}$
 $B_7 = \{2, 3, 5\}$ $B_8 = \{2, 3\}$

Item #	Count
1	5
2	5
3	5
4	2
5	4

The hash table is

Bucket #	Count
0	6
1	2
2	4
3	4
4	5

For $s=3$, $L_1 = \{1, 2, 3, 5\}$, and Bitmap $\{1, 0, 1, 1, 1\}$

Observations about Buckets

The hash table is

Bucket #	Count
0	6
1	2
2	4
3	4
4	5

For $s=3$, $L_1 = \{1, 2, 3, 5\}$

- If the count of a bucket is \geq support s , it is called a **frequent bucket**
- For a bucket with total count less than s , none of its pairs can be frequent. Can be eliminated as candidates!
- For Pass 2, only count pairs that hash to frequent buckets

PCY Algorithm - Pass 2

- Count all pairs $\{i, j\}$ that meet the conditions
 1. Both i and j are frequent items
 2. The pair $\{i, j\}$ hashed to a frequent bucket
(count $\geq s$)
- All these conditions are necessary for the pair to have a chance of being frequent

PCY Algorithm - Pass 2

Hash table after pass 1:

The hash table is

Bucket #	Count
0	6
1	2
2	4
3	4
4	5

For $s=3$, $L_1 = \{1, 2, 3, 5\}$, and Bitmap $\{1, 0, 1, 1, 1\}$

Pass 2

Frequent items are $\{1, 2, 3, 5\}$

Candidate pairs and their counts

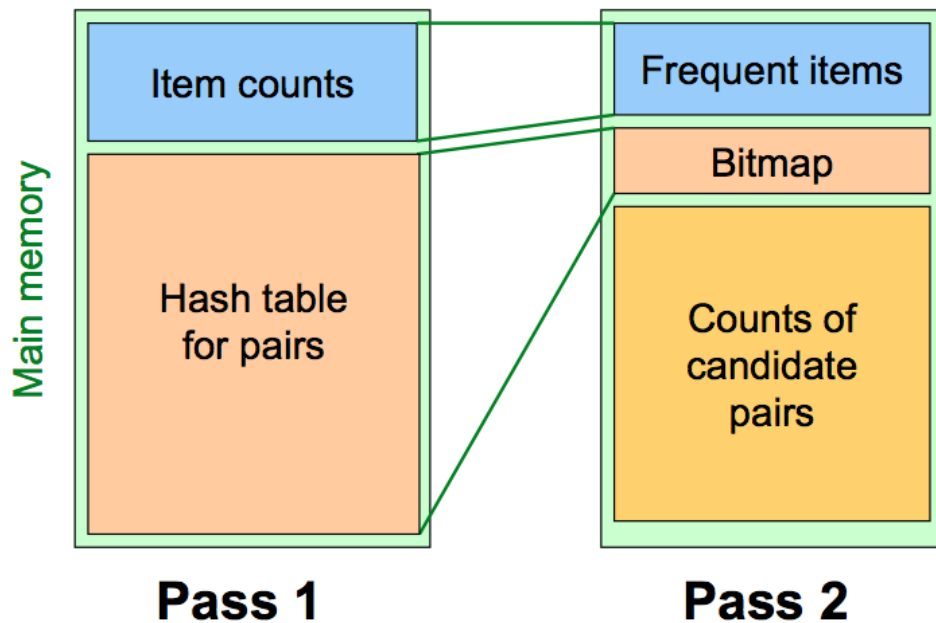
Pair	Count
(2,3)	4
(2,5)	3
(1,2)	2
(3,5)	2
(1,3)	4

$(1, 4), (2, 3) \rightarrow h(i, j) = 0$
 $(1, 5), (2, 4) \rightarrow h(i, j) = 1$
 $(2, 5), (3, 4) \rightarrow h(i, j) = 2$
 $(1, 2), (3, 5) \rightarrow h(i, j) = 3$
 $(1, 3), (4, 5) \rightarrow h(i, j) = 4$

frequent itemsets are

$\{1\}, \{2\}, \{3\}, \{5\}, \{1, 3\}, \{2, 3\}, \{2, 5\}$

Main-Memory: Picture of PCY

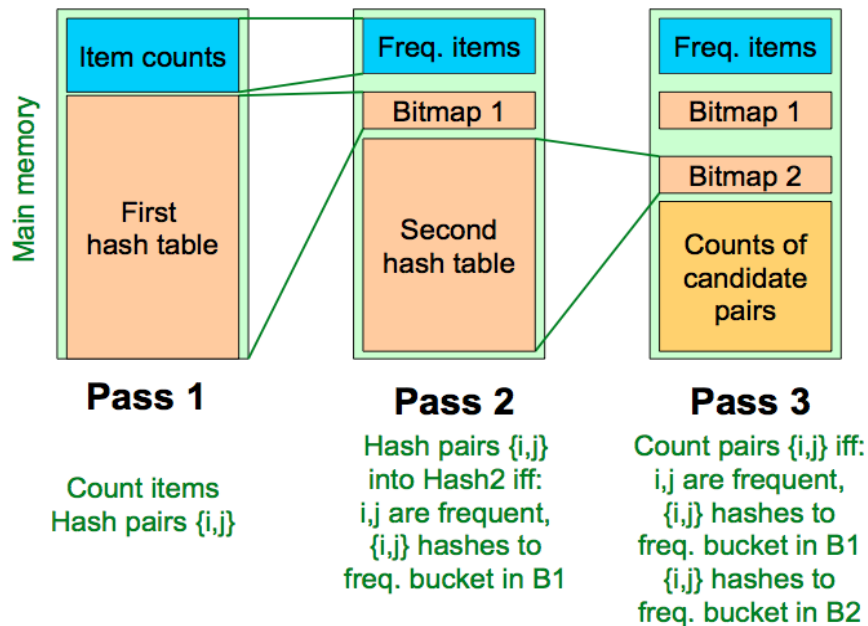


Refinement

- Remember: Memory is the **bottleneck!**
- Can we further limit the number of candidates to be counted?
- Refinement for PCY Algorithm
 - Multistage
 - Multihash

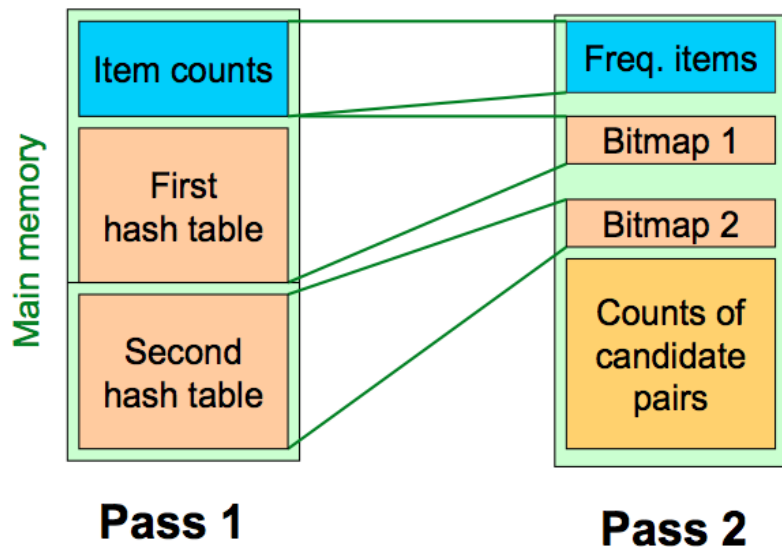
Multistage Algorithm

- **Key Idea:** After Pass 1 of PCY, rehash only those pairs that qualify for pass 2 of PCY
- Require additional pass over the data
- Important points
 - Two hash functions have to be independent
 - Check both hashes on the third pass



Multihash Algorithm

- **Key Idea:** Use several independent hash functions on the first pass
- **Risk:** Halving the number of buckets **doubles** the average count
- If most buckets still not reach count s , then we can get a benefit like multistage, but in only 2 passes!
- Possible candidate pairs $\{i, j\}$:
 - i, j are frequent items
 - $\{i, j\}$ are hashed into both frequent buckets




Frequent Itemsets in ≤ 2 Passes

- A-Priori, PCY, etc., take k passes to find frequent itemsets of size k
- Can we use fewer passes?
- Use 2 or fewer passes for all sizes
 - Random sampling
 - may miss some frequent itemsets
 - SON (Savasere, Omiecinski, and Navathe)
 - Toivonen (not going to conver)

Random Sampling

- Take a random sample of the market baskets
- Run A-priori in **main memory**
 - Don't have to pay for disk I/O each time we read over the data
 - Reduce the support threshold proportionally to match the sample size (e.g. 1% of Data, support $\Rightarrow 1/100*s$)
- Verify the candidate pairs by a second pass



SON Algorithm



Chunks of baskets



Run a-priori with
 $(1/n) \cdot \text{support}$

Memory

Save all the possible
candidates of each
chunk




Disk

- Repeatedly read small subsets of the baskets into main memory and run an in-memory algorithm to find all frequent itemsets
- Possible candidates:
 - Union all the frequent itemsets found in each chunk
 - why? “**monotonicity**” idea: an itemset cannot be frequent in the entire set of baskets unless it is frequent in at least one subset
- On a **second pass**, count all the candidate

SON Algorithm- Distributed Version

● MapReduce for Pass 1

Chunks of baskets



- Distributed data mining
- Pass 1: Find candidate itemsets
 - Map: $(F, 1)$
 - F : frequent itemset
 - Reduce: Union all the $(F, 1)$
- Pass 2: Find true frequent itemsets
 - Map: (C, v)
 - C : possible candidate
 - Reduce: Add all the (C, v)

FP-Growth Approach

Introduction

- A-priori
 - Generation of candidate itemset (Expensive in both space and time)
 - Support counting is expensive
 - Subset checking
 - Multiple Database scans (I/O)

FP-Growth approach


- FP-Growth (Frequent Pattern-Growth)
 - Mining in main memory to reduce (#DBscans)
 - Without candidate itemsets generation
- Two step approach
 - Step 1: Build a compact data structure called the FP-tree
 - Step 2: Extracts frequent itemsets directly from the FP-tree (Traversal through FP-tree)

FP-Tree construction

- FP-Tree construction
 - Pass 1:
 - Find the frequent items
 - Pass 2:
 - Construct FP-Tree

f	4
c	4
a	3
b	3
m	3
p	3


TID	Items bought	Ordered
100	{a, c, d, f, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, i, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, c, e, f, l, m, n, p}	{f, c, a, m, p}



FP-Tree

- FP-Tree
 - Prefix Tree
 - Has a much smaller size than the uncompressed data
 - Mining in main memory
- How to find the Frequent itemset?
 - Tree traversal
 - Bottom-up algorithm
 - Divide and conquer
 - More detail:

<http://csc.lsu.edu/~jianhua/FPGrowth.pdf>



FP-Growth V.S A-priori

	Apriori	FP-Growth
# Passes over data	depends	2
Candidate Generation	Yes	No

- FP-Growth Pros:
 - “Compresses” data-set, mining in memory
 - much faster than Apriori
- FP-Growth Cons:
 - FP-Tree may not fit in memory
 - FP-Tree is expensive to build
 - Trade-off: takes time to build, but once it is build, frequent itemsets are read off easily

Acknowledgements

- Stanford CS246: Mining Massive Datasets (Jure Leskovec)
- Mining of Massive Datasets (Anand Rajaraman, Jeffrey Ullman)
- Introduction to Frequent Pattern Growth (FP-Growth) Algorithm (Florian Verhein)
- NCCU: Data-mining (Man-Kwan Shan)
- [Mining frequent patterns without candidate generation. A frequent-tree approach](#), SIGMOD '00 Proceedings of the 2000